

Amazon Product Reviews

July 16, 2020

1 Amazon Product Reviews Dataset

<https://www.kaggle.com/saurav9786/amazon-product-reviews>

userId : Every user identified with a unique id

productId : Every product identified with a unique id

Rating : Rating of the corresponding product by the corresponding user

timestamp : Time of the rating (ignore this column for this exercise)

```
[1]: import pandas as pd
import numpy as np
import os
import math
import json
import time
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.model_selection import train_test_split
from sklearn.neighbors import NearestNeighbors

import scipy.sparse
from scipy.sparse import csr_matrix
from scipy.sparse.linalg import svds

import warnings; warnings.simplefilter('ignore')
%matplotlib inline
```

```
[2]: # Read and give names to data columns
data = pd.read_csv('../data/ratings_electronics.csv',
                  names = ['userId', 'productId', 'Rating', 'timestamp'])
```

1.1 Data Exploration and Visualization

```
[3]: data.head()
```

```
[3]:      userId  productId  Rating  timestamp
0  AKM1MP6P00YPR  0132793040    5.0  1365811200
1  A2CX7LUOHB2NDG  0321732944    5.0  1341100800
2  A2NWSAGRHC8P8N5  0439886341    1.0  1367193600
3  A2WNBOD3WNDNKT  0439886341    3.0  1374451200
4  A1GI0U4ZRJA8WN  0439886341    1.0  1334707200
```

```
[4]: print('Data shape: ', data.shape)
```

Data shape: (7824482, 4)

Use only 1,000,000 datapoints

```
[5]: data = data.iloc[:1000000,:]
```

```
[6]: print('Data shape: ', data.shape)
```

Data shape: (1000000, 4)

```
[7]: print('Data types:\n',data.dtypes)
```

Data types:

userId	object
productId	object
Rating	float64
timestamp	int64
dtype:	object

```
[8]: print("Describe: ", data.describe()['Rating'])
```

Describe: count 1000000.000000

mean	3.973620
std	1.399741
min	1.000000
25%	3.000000
50%	5.000000
75%	5.000000
max	5.000000

Name: Rating, dtype: float64

```
[9]: min_rating, max_rating = data.Rating.min(), data.Rating.max()
print('Minimum: {} and maximum: {} rating '.format(min_rating, max_rating ))
```


Minimum: 1.0 and maximum: 5.0 rating

1.2 Data Cleaning

```
[10]: print("Number of missing values:",  
        data.isnull().sum()  
        )
```

```
Number of missing values: userId      0  
productId    0  
Rating       0  
timestamp    0  
dtype: int64
```

```
[11]: with sns.axes_style('white'):  
        g = sns.factorplot("Rating", data=data, aspect=2.0, kind='count')  
        g.set_ylabels("Total number of ratings")
```



```
[12]: print("Some statistics ")  
        print("-"*50)  
        print("\nTotal no of ratings :", data.shape[0])  
        print("Total No of Users   :", len(np.unique(data.userId)))  
        print("Total No of products  :", len(np.unique(data.productId)))
```

Some statistics

```
-----  
  
Total no of ratings : 1000000  
Total No of Users   : 754153  
Total No of products : 59634
```

Drop timestamp column

```
[13]: data.drop(['timestamp'], axis = 1, inplace = True)
```

```
[14]: data.head()
```

```
[14]:
```

	userId	productId	Rating
0	AKM1MP6P00YPR	0132793040	5.0
1	A2CX7LUOHB2NDG	0321732944	5.0
2	A2NWSAGRHC8P8N5	0439886341	1.0
3	A2WNBOD3WNDNKT	0439886341	3.0
4	A1GIOU4ZRJA8WN	0439886341	1.0

```
[15]: data.columns
```

```
[15]: Index(['userId', 'productId', 'Rating'], dtype='object')
```

Only products with 100 or more ratings

```
[16]: data_count_ratings = data.groupby('productId').filter(lambda d: d['Rating'].  
    ↳ count() >= 100 )
```

```
[17]: selected_items = set(data_count_ratings['productId'])  
print("Len of selected items: ", len(selected_items))  
print(list(selected_items)[:10])
```

```
Len of selected items: 1712  
['B00005BMSN', 'B00009705F', 'B0002ZQHFA', 'B00083Y0YG', 'B00009XVA3',  
'B00070WNCC', 'B00004TWM6', 'B000CBB4N4', 'B00004Z61H', 'B000FED6N0']
```

```
[18]: data.shape
```

```
[18]: (1000000, 3)
```

```
[19]: ### Only products with 100 or more reviews  
data = data[data.productId.isin(selected_items)]
```

```
[20]: data.shape
```

```
[20]: (509370, 3)
```

```
[21]: data.iloc[15]
```

```
[21]:
```

userId	A1ZD73MDX4POAY
productId	0972683275
Rating	5

```
Name: 198, dtype: object
```

1.2.1 Training and testing sets

```
[22]: n,m = data.shape

training_percentage = 0.8
n_training = int(n*0.8)

training, test = data.iloc[:n_training], data.iloc[n_training:]
print( 'Length training {} and test {}'.format(len(training), len(test)))
```

Length training 407496 and test 101874

2 Recommender System

```
[23]: from collections import defaultdict

def create_sets(data):
    """
    data:Pandas dataframe
    """
    n,m = data.shape

    items_per_user = defaultdict(set)
    users_per_item = defaultdict(set)

    for i in range(n):
        datapoint = data.iloc[i]
        user,item = datapoint['userId'], datapoint['productId']
        items_per_user[user].add(item)
        users_per_item[item].add(user)

    return items_per_user ,users_per_item

def create_set_reviews(data):
    reviewsPerUser = defaultdict(list)
    reviewsPerItem = defaultdict(list)

    n,m = data.shape

    for j in range(n):
        d = data.iloc[j]
        user,item = d['userId'], d['productId']
        reviewsPerUser[user].append(d)
        reviewsPerItem[item].append(d)
```

```
return reviewsPerUser, reviewsPerItem
```

```
[24]: items_per_user ,users_per_item = create_sets(training)
```

```
[25]: reviewsPerUser, reviewsPerItem = create_set_reviews(training)
```

```
[28]: n,m = training.shape

overall_mean_rating = training['Rating'].mean()
overall_mean_rating
```

```
[28]: 4.096099102813279
```

```
[29]: def Jaccard(s1, s2):
    numer = len(s1.intersection(s2))
    denom = len(s1.union(s2))
    return numer / denom

def mostSimilar(iD, n, reviewsPerUser, usersPerItem):
    similarities = []
    users = usersPerItem[iD]
    for i2 in usersPerItem:
        if i2 == iD: continue
        sim = Jaccard(users, usersPerItem[i2])
        similarities.append((sim,i2))
    similarities.sort(reverse=True)
    return similarities[:n]
```

```
[43]: def predictRating(user, item, reviewsPerUser, usersPerItem):
    ratings = []
    similarities = []

    for d in reviewsPerUser[user]:

        # product
        i2 = d['productId']

        if i2 == item: continue

        ratings.append(d['Rating'])

        similarities.append(Jaccard(usersPerItem[item], usersPerItem[i2]))
```

```

if (sum(similarities) > 0):
    weightedRatings = [(x*y) for x,y in zip(ratings,similarities)]
    return sum(weightedRatings) / sum(similarities)

else:
    # User hasn't rated any similar items
    return overall_mean_rating

```

2.1 Make Recommendation

```

[44]: item = training.iloc[200]['productId']
      item

```

```

[44]: '0972683275'

```

```

[45]: mostSimilar(iD = item, n= 10 ,reviewsPerUser=reviewsPerUser, usersPerItem=
      ↪=users_per_item )

```

```

[45]: [(0.0020174848688634837, 'B00010HH0Q'),
      (0.0017391304347826088, 'B0002855KK'),
      (0.001729106628242075, 'B00005ML7Q'),
      (0.001594896331738437, 'B0002GV876'),
      (0.0015936254980079682, 'B000A2AGYS'),
      (0.0015923566878980893, 'B0007A1IRC'),
      (0.0014803849000740192, 'B0006I09LQ'),
      (0.0013708019191226869, 'B000ARAPQW'),
      (0.0013201320132013201, 'B00005T3N3'),
      (0.0013192612137203166, 'B0007MWE1E')]

```

2.1.1 Evaluate Performace

- Compare against mean rating by user

```

[46]: def MSE(predictions, labels):
      differences = [(x-y)**2 for x,y in zip(predictions,labels)]
      return sum(differences) / len(differences)

```

```

[47]: # Overall mean rating
      n,m = training.shape
      mean_rating = [overall_mean_rating]*n
      n,m = training.shape
      labels = [training.iloc[k]['Rating'] for k in range(n)]

```

```

[50]: mean_rating[:10]

```

[illegible]

```
[48]: labels[:10]
```

```
[48]: [4.0, 4.0, 5.0, 4.0, 5.0, 4.0, 5.0, 3.0, 5.0, 5.0]
```

```
[52]: predictions = []

for i in range(n):
    d = training.iloc[i]
    user = d['userId']
    item = d['productId']

    stars = predictRating(user, item, reviewsPerUser=reviewsPerUser,
↪ usersPerItem = users_per_item)
    predictions.append(stars)
```

```
[53]: predictions[50:100]
```

[illegible]


```

4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
5.0,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.0,
4.096099102813279,
5.0,
4.096099102813279,
4.096099102813279,
4.0,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279,
4.096099102813279]

```

```
[54]: print(MSE(mean_rating, labels), MSE(predictions, labels))
```

```
1.7587234209338283 1.917202902116438
```

2.2 Conclusion

Collaborative filtering method not always provide the best solution, weighting previous ratings can be or not be a good prediction for future ratings.

2.3 Model-based collaborative filtering system

Item-user matrix M $M[i,j]$ \$ rating given to item i by user j

```
[ ]:
```

```

def create_item_user_matrix(data):
    M = data.pivot_table(values = 'Rating', index = 'userId', columns = '
    ↪productId', fill_value = 0)
    M = M.T

    return M

M = create_item_user_matrix(training)

M.head()

from sklearn.decomposition import TruncatedSVD

def correlation(M, n_componets):
    """
    M: matrix of items-products
    ratings

    return correlation matrix tem-item
    most similar items are more correlated
    """

    # Decompose
    SVD = TruncatedSVD(n_components = 10 )
    decomposed = SVD.fit_transform(M)
    correlation = np.corrcoef(decomposed)

    return correlation

correlation = correlation(M, n_componets= 10 )
items_names = list(M.index)

i = 100
item_name = M.index[i]
print('Item name: ', item_name)
item_index = items_names.index(item_name)
print('Item index: ', item_index)

### Recommend items based on item B000021YU8

```

```

# Correlations for item B000021YU8
correlations_item = correlation[item_index]
print("Correlations shape {} for item {}".format(correlations_item.shape,
→item_name))

# recommend items with correlation > 0.65
r_items = list(M.index[correlations_item > 0.65] )

# Remove the item itself
r_items.remove(item_name)

print('Recommended first best 24 items: {}, based on item {}'.format(r_items,
→item_name))

# Correlations for item B000021YU8
correlations_item = correlation[item_index]
print("Correlations shape {} for item {}".format(correlations_item.shape,
→item_name))

# recommend items with correlation > 0.65
r_items = list(M.index[correlations_item > 0.65] )

# Remove the item itself
r_items.remove(item_name)

print('Recommended first best 24 items: {}, based on item {}'.format(r_items,
→item_name))

```

This method uses item-item correlations to recommend more products to an user. It decompose the rating matrix M , representing each item as a vector (embedding).

The method works well with large and sparse matrices. Allow us to recommend many items to users, but with limited personalization.

This approach ignores information user-user about ratings.

[]: